

Advanced Machine Learning

LAB 3: Reinforcement Learning

Mohammed Bakheet (mohba508)

07/10/2020

Contents

Question 1: Q-Learning	2
Greedy Policy Function	2
Epsilon Greedy Function	2
Transition Model	3
Q Learning Function	4
Environment A:	4
Answers:	4
Environment B:	8
Answer:	8
Environment C	19
Answer:	19
Question 2: Reinforce:	23
Visualizing the environment:	23
Transition Model:	23
Deep Policy Distribution:	23
Deep Policy	24
Deep Policy Training	24
Reinforcement Episode	25
Environment D:	26
Answer:	26
Environment E:	45
Answer:	45
References	58
Appendix	58

Question 1: Q-Learning

The file RL Lab1.R in the course website contains a template of the Q-learning algorithm. 1 You are asked to complete the implementation. We will work with a grid-world environment consisting of $H \times W$ tiles laid out in a 2-dimensional grid. An agent acts by moving up, down, left or right in the grid-world. This corresponds to the following Markov decision process:

- State space: $S = \{(x, y) \mid x \in \{1, \dots, H\}, y \in \{1, \dots, W\}\}$.
- Action space: $A = \{\text{up, down, left, right}\}$.

Additionally, we assume state space to be fully observable. The reward function is a deterministic function of the state and does not depend on the actions taken by the agent. We assume the agent gets the reward as soon as it moves to a state. The transition model is defined by the agent moving in the direction chosen with probability $(1 - \beta)$. The agent might also slip and end up moving in the direction to the left or right of its chosen action, each with probability $\beta/2$. The transition model is unknown to the agent, forcing us to resort to model-free solutions. The environment is episodic and all states with a non-zero reward are terminal. Throughout this lab we use integer representations of the different actions: Up=1, right=2, down=3 and left=4.

Greedy Policy Function

```
GreedyPolicy <- function(x, y){  
  
  # Get a greedy action for state (x,y) from q_table.  
  #  
  # Args:  
  #   x, y: state coordinates.  
  #   q_table (global variable): a HxWx4 array containing Q-values for each state-action pair.  
  #  
  # Returns:  
  #   An action, i.e. integer in {1,2,3,4}.  
  
  # Your code here.  
  rewards = q_table[x,y,1:4]  
  action = 0  
  max_reward = which(rewards == max(rewards))  
  
  if (length(max_reward)>1){  
    action = sample(max_reward,1)  
  }  
  else  
  {  
    action = max_reward  
  }  
  return(action)  
}
```

Epsilon Greedy Function

```
EpsilonGreedyPolicy <- function(x, y, epsilon){  
  
  # Get an epsilon-greedy action for state (x,y) from q_table.  
  #
```

```

# Args:
#   x, y: state coordinates.
#   epsilon: probability of acting greedily.
#
# Returns:
#   An action, i.e. integer in {1,2,3,4}.

# Your code here.
actions = c(1:4)
action = 0
rewards = q_table[x,y,1:4]

max_reward = which(rewards == max(rewards))

if (length(max_reward)>1){
  action = sample(max_reward,1)
}
else
{
  action = max_reward
}

if (1-epsilon < runif(1)){
  return(sample(actions,1))
}
else{
  return(action)
}
}

```

Transition Model

```

transition_model <- function(x, y, action, beta){

  # Computes the new state after given action is taken. The agent will follow the action
  # with probability (1-beta) and slip to the right or left with probability beta/2 each.
  #
  # Args:
  #   x, y: state coordinates.
  #   action: which action the agent takes (in {1,2,3,4}).
  #   beta: probability of the agent slipping to the side when trying to move.
  #   H, W (global variables): environment dimensions.
  #
  # Returns:
  #   The new state after the action has been taken.

  delta <- sample(-1:1, size = 1, prob = c(0.5*beta,1-beta,0.5*beta))
  final_action <- ((action + delta + 3) %% 4) + 1
  foo <- c(x,y) + unlist(action_deltas[final_action])
  foo <- pmax(c(1,1),pmin(foo,c(H,W)))

  return (foo)
}

```

Q Learning Function

Environment A:

For our first environment, we will use $H = 5$ and $W = 7$. This environment includes a reward of 10 in state (3,6) and a reward of -1 in states (2,3), (3,3) and (4,3).

We specify the rewards using a reward map in the form of a matrix with one entry for each state. States with no reward will simply have a matrix entry of 0. The agent starts each episode in the state (3,1). The function `vis_environment` in the file `RL Lab1.R` is used to visualize the environment and learned action values and policy. You will not have to modify this function, but read the comments in it to familiarize with how it can be used.

When implementing Q-learning, the estimated values of $Q(S, A)$ are commonly stored in a data-structured called Q-table. This is nothing but a tensor with one entry for each state-action pair. Since we have a $H \times W$ environment with four actions, we can use a 3D-tensor of dimensions $H \times W \times 4$ to represent our Q-table. Initialize all Q-values to 0. Run the function `vis_environment` before proceeding further. Note that each non-terminal tile has four values. These represent the action values associated to the tile (state). Note also that each non-terminal tile has an arrow. This indicates the greedy policy for the tile (ties are broken at random).

You are requested to carry out the following tasks:

- Implement the greedy and ϵ -greedy policies in the functions `GreedyPolicy` and `EpsilonGreedyPolicy` of the file `RL Lab1.R`. The functions should break ties at random, i.e. they should sample uniformly from the set of actions with maximal Q-value.

- Implement the Q-learning algorithm in the function `q_learning` of the file `RL Lab1.R`.

The function should run one episode of the agent acting in the environment and update the Q-table accordingly. The function should return the episode reward and the sum of the temporal-difference correction terms $R + \gamma * \max_a Q(S', a) - Q(S, A)$ for all steps in the episode. Note that a transition model taking β as input is already implemented for you in the function `transition model`.

- Run 10000 episodes of Q-learning with $\epsilon = 0.5$, $\beta = 0$, $\alpha = 0.1$ and $\gamma = 0.95$. To do so, simply run the code provided in the file `RL Lab1.R`. The code visualizes the Q-table and a greedy policy derived from it after episodes 10, 100, 1000 and 10000. Answer the following questions:

- What has the agent learned after the first 10 episodes ?
 - Is the final greedy policy (after 10000 episodes) optimal? Why / Why not ?
 - Does the agent learn that there are multiple paths to get to the positive reward ?
- If not, what could be done to make the agent learn this ?

Answers:

- After the first 10 episodes the agent hasn't learned much, it's got knowledge about avoiding the states with negative ones (-1).
- After 10000 iterations, the agent has learned how to reach the goal following the optimal path, in some states i.e (1,3) the goodness of some actions was the same, but still the agent was able to pick the correct path to the target state. if we increase the number of iterations the accuracy will increase, but we are satisfied with the current policy and we can say the solution is optimal.
- The agent has learned how to reach to the goal from the state following one paths, it takes one path to the goal and it doesn't slip due the β value of 0 that doesn't allow the agent to slip from the path. And if we want the agent to try different paths, then we can increase the value of β that allows the agent to slip from the path and try different paths to the target.

```
# Environment A (learning)
```

```
H <- 5
```

```

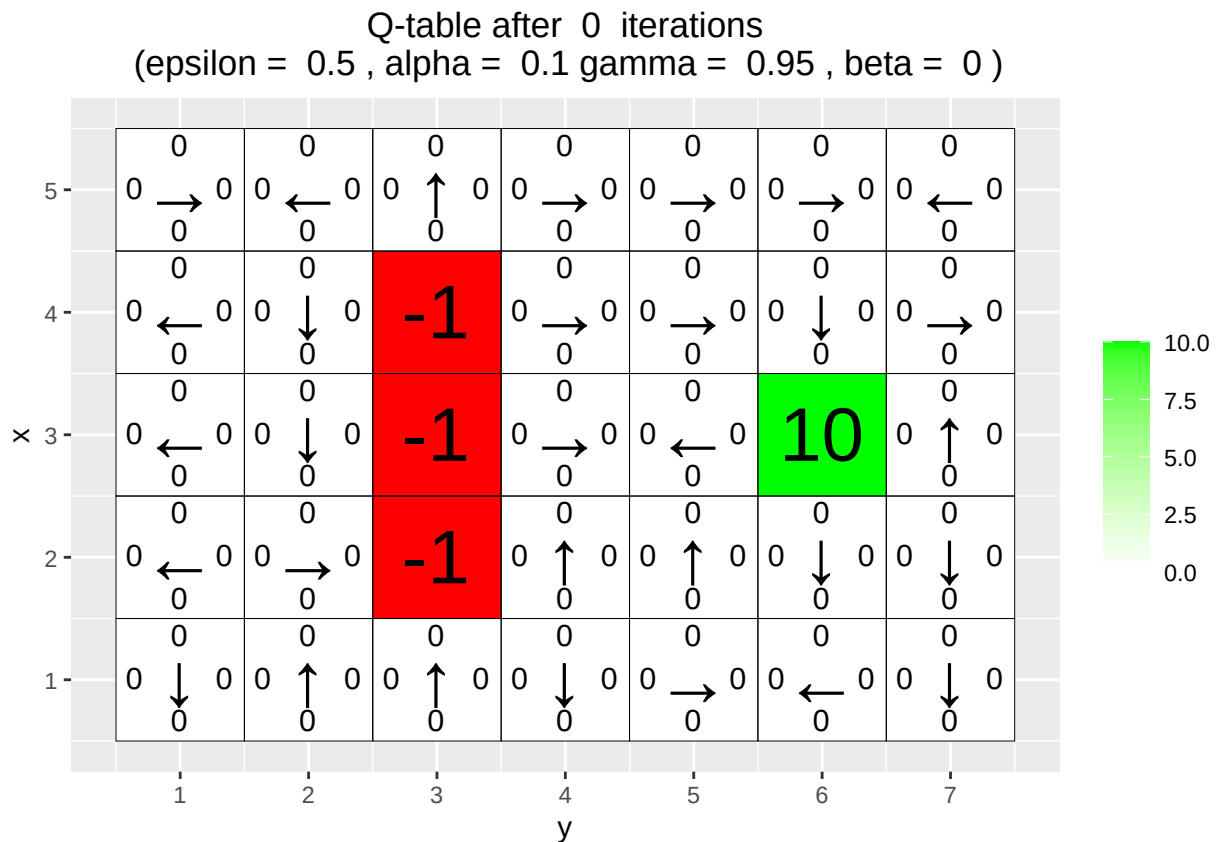
W <- 7

reward_map <- matrix(0, nrow = H, ncol = W)
reward_map[3,6] <- 10
reward_map[2:4,3] <- -1

q_table <- array(0,dim = c(H,W,4))

vis_environment()

```




```

for(i in 1:10000){
  foo <- q_learning(start_state = c(3,1))

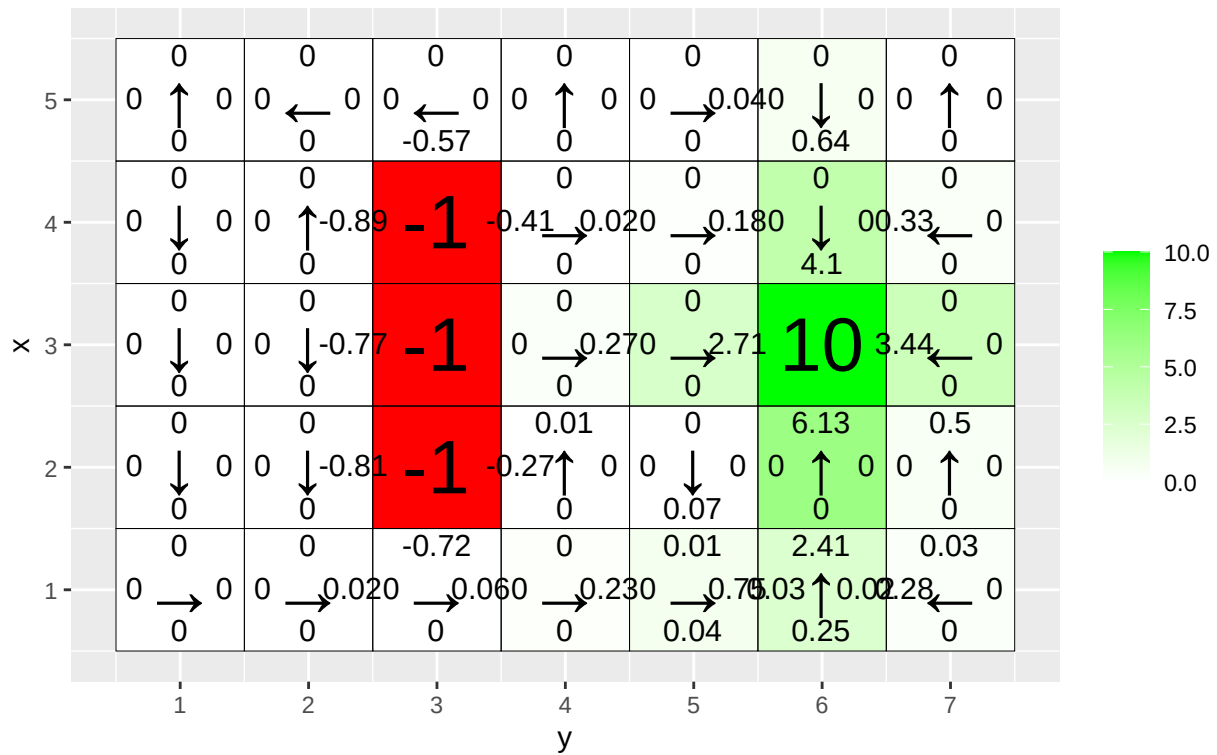
  if(any(i==c(10,100,1000,10000)))
    vis_environment(i)
}

```


Q-table after 10 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.95 , beta = 0)




Q-table after 100 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.95 , beta = 0)



Q-table after 1000 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.95 , beta = 0)



Q-table after 10000 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.95 , beta = 0)



Environment B:

This is a 7×8 environment where the top and bottom rows have negative rewards. In this environment, the agent starts each episode in the state (4, 1). There are two positive rewards, of 5 and 10. The reward of 5 is easily reachable, but the agent has to navigate around the first reward in order to find the reward worth 10. Your task is to investigate how the ϵ and γ parameters affect the learned policy by running 30000 episodes of Q-learning with $\epsilon = 0.1, 0.5, \gamma = 0.5, 0.75, 0.95, \beta = 0$ and $\alpha = 0.1$. To do so, simply run the code provided in the file RL Lab1.R and explain your observations.

Answer:

As we can see from the plot when:

($\epsilon = 0.1, \alpha = 0.1, \gamma = 0.5, \beta = 0$), vs ($\epsilon = 0.1, \alpha = 0.1, \gamma = 0.75, \beta = 0$) vs ($\epsilon = 0.1, \alpha = 0.1, \gamma = 0.95, \beta = 0$)

When we increased the value of γ from 0.5 to 0.75 with the same value of ϵ , the agent hasn't explored after reaching the state with 5 reward because we increased the discount factor that determines whether the agent should take the immediate reward or explore (considers rewards later in time). And when using an γ value of 0.95, the agent has learned more about how to reach the state with 5 rewards (if we have the same preferences, then γ should be 1), but still the agent has no idea about the state with 10 rewards because the value of ϵ is small which means not taking a greedy action exploring with the discount factor, and the number of iterations also effect the learning γ

Nonetheless, when the value of ϵ increases to 5, then the agent learns about the state with 10 reward, and that is because a higher value of ϵ allows the agent to explore the environment more by not taking the action that maximizes the reward all the time, but with probability.

```
# Environment B (the effect of epsilon and gamma)
```

```
H <- 7
```

```
W <- 8
```

```
reward_map <- matrix(0, nrow = H, ncol = W)
```

```
reward_map[1,] <- -1
```

```
reward_map[7,] <- -1
```


```
reward_map[4,5] <- 5
```

```
reward_map[4,8] <- 10
```

```
q_table <- array(0,dim = c(H,W,4))
```

```
vis_environment()
```


Q-table after 0 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.95 , beta = 0)



```
MovingAverage <- function(x, n){
  cx <- c(0,cumsum(x))
  rsum <- (cx[(n+1):length(cx)] - cx[1:(length(cx) - n)]) / n



  return (rsum)
}


for(j in c(0.5,0.75,0.95)){
  q_table <- array(0,dim = c(H,W,4))
  reward <- NULL
  correction <- NULL

  for(i in 1:30000){
    foo <- q_learning(gamma = j, start_state = c(4,1))
    reward <- c(reward,foo[1])
    correction <- c(correction,foo[2])
  }


  vis_environment(i, gamma = j)
  plot(MovingAverage(reward,100),type = "l")
  plot(MovingAverage(correction,100),type = "l")
}
```



Q-table after 30000 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.5 , beta = 0)








Index
Q-table after 30000 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.75 , beta = 0)





Q-table after 30000 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.95 , beta = 0)








```
for(j in c(0.5,0.75,0.95)){
  q_table <- array(0,dim = c(H,W,4))
  reward <- NULL
  correction <- NULL

  for(i in 1:30000){
    foo <- q_learning(epsilon = 0.1, gamma = j, start_state = c(4,1))
    reward <- c(reward,foo[1])
    correction <- c(correction,foo[2])
  }


  vis_environment(i, epsilon = 0.1, gamma = j)
  plot(MovingAverage(reward,100),type = "l")
  plot(MovingAverage(correction,100),type = "l")
}
```



Q-table after 30000 iterations
(epsilon = 0.1 , alpha = 0.1 gamma = 0.5 , beta = 0)








Index
Q-table after 30000 iterations
(epsilon = 0.1 , alpha = 0.1 gamma = 0.75 , beta = 0)





Q-table after 30000 iterations
(epsilon = 0.1 , alpha = 0.1 gamma = 0.95 , beta = 0)





Environment C

This is a smaller 3×6 environment. Here the agent starts each episode in the state (1,1). Your task is to investigate how the β parameter affects the learned policy by running 10000 episodes of Q-learning with $\beta = 0, 0.2, 0.4, 0.66$, $\epsilon = 0.5$, $\gamma = 0.6$ and $\alpha = 0.1$. To do so, simply run the code provided in the file RL Lab1.R and explain your observations.

Answer:

As it's seen in the plots when we increase the value of β , the frequency of the agent slipping from the path to the goal increases, for the highest values of β (0.66) it goes in a loop in some states.

```
# Environment C (the effect of beta).


H <- 3
W <- 6

reward_map <- matrix(0, nrow = H, ncol = W)
reward_map[1,2:5] <- -1
reward_map[1,6] <- 10

q_table <- array(0,dim = c(H,W,4))

vis_environment()
```

Q-table after 0 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.95 , beta = 0)




```
for(j in c(0,0.2,0.4,0.66)){
  q_table <- array(0,dim = c(H,W,4))


  for(i in 1:10000)
    foo <- q_learning(gamma = 0.6, beta = j, start_state = c(1,1))

  vis_environment(i, gamma = 0.6, beta = j)
}
```


Q-table after 10000 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.6 , beta = 0)




Q-table after 10000 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.6 , beta = 0.2)



Q-table after 10000 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.6 , beta = 0.4)



Q-table after 10000 iterations
(epsilon = 0.5 , alpha = 0.1 gamma = 0.6 , beta = 0.66)



Question 2: Reinforce:

The file RL Lab2.R in the course website contains an implementation of the REINFORCE algorithm. 2 Your task is to run the code provided and answer some questions. Although you do not have to modify the code, you are advised to check it out to familiarize with it. The code uses the R package keras to manipulate neural networks.

We will work with a 4×4 grid. We want the agent to learn to navigate to a random goal position in the grid. The agent will start in a random position and it will be told the goal position.

The agent receives a reward of 5 when it reaches the goal. Since the goal position can be any position, we need a way to tell the agent where the goal is. Since our agent does not have any memory mechanism, we provide the goal coordinates as part of the state at every time step, i.e. a state consists now of four coordinates: Two for the position of the agent, and two for the goal position. The actions of the agent can however only impact its own position, i.e. the actions do not modify the goal position. Note that the agent initially does not know that the last two coordinates of a state indicate the position with maximal reward, i.e. the goal position. It has to learn it. It also has to learn a policy to reach the goal position from the initial position. Moreover, the policy has to depend on the goal position, because it is chosen at random in each episode. Since we only have a single non-zero reward, we do not specify a reward map. Instead, the goal coordinates are passed to the functions that need to access the reward function.

Visualizing the environment:

Transition Model:

```
transition_model <- function(x, y, action, beta){  
  
  # Computes the new state after given action is taken. The agent will follow the action  
  # with probability (1-beta) and slip to the right or left with probability beta/2 each.  
  #  
  # Args:  
  #   x, y: state coordinates.  
  #   action: which action the agent takes (in {1,2,3,4}).  
  #   beta: probability of the agent slipping to the side when trying to move.  
  #   H, W (global variables): environment dimensions.  
  #  
  # Returns:  
  #   The new state after the action has been taken.  
  
  delta <- sample(-1:1, size = 1, prob = c(0.5*beta,1-beta,0.5*beta))  
  final_action <- ((action + delta + 3) %% 4) + 1  
  foo <- c(x,y) + unlist(action_deltas[final_action])  
  foo <- pmax(c(1,1),pmin(foo,c(H,W)))  
  
  return (foo)  
}
```

Deep Policy Distribution:

```
DeepPolicy_dist <- function(x, y, goal_x, goal_y){  
  
  # Get distribution over actions for state (x,y) and goal (goal_x,goal_y) from the deep policy.  
  #  
  # Args:  
  #   x, y: state coordinates.
```

```

#   goal_x, goal_y: goal coordinates.
#   model (global variable): NN encoding the policy.
#
# Returns:
#   A distribution over actions.

foo <- matrix(data = c(x,y,goal_x,goal_y), nrow = 1)

# return (predict_proba(model, x = foo))
return (predict_on_batch(model, x = foo)) # Faster.
}

```

Deep Policy

```

DeepPolicy <- function(x, y, goal_x, goal_y){

  # Get an action for state (x,y) and goal (goal_x,goal_y) from the deep policy.
  #
  # Args:
  #   x, y: state coordinates.
  #   goal_x, goal_y: goal coordinates.
  #   model (global variable): NN encoding the policy.
  #
  # Returns:
  #   An action, i.e. integer in {1,2,3,4}.

  foo <- DeepPolicy_dist(x,y,goal_x,goal_y)

  return (sample(1:4, size = 1, prob = foo))
}

```

Deep Policy Training

```

DeepPolicy_train <- function(states, actions, goal, gamma){

  # Train the policy network on a rolled out trajectory.
  #
  # Args:
  #   states: array of states visited throughout the trajectory.
  #   actions: array of actions taken throughout the trajectory.
  #   goal: goal coordinates, array with 2 entries.
  #   gamma: discount factor.

  # Construct batch for training.
  inputs <- matrix(data = states, ncol = 2, byrow = TRUE)
  inputs <- cbind(inputs,rep(goal[1],nrow(inputs)))
  inputs <- cbind(inputs,rep(goal[2],nrow(inputs)))

  targets <- array(data = actions, dim = nrow(inputs))
  targets <- to_categorical(targets-1, num_classes = 4)
}

```



```

# Sample weights. Reward of 5 for reaching the goal.
weights <- array(data = 5*(gamma^(nrow(inputs)-1)), dim = nrow(inputs))

# Train on batch. Note that this runs a SINGLE gradient update.
train_on_batch(model, x = inputs, y = targets, sample_weight = weights)
}

```

Reinforcement Episode

```

reinforce_episode <- function(goal, gamma = 0.95, beta = 0){

  # Rolls out a trajectory in the environment until the goal is reached.
  # Then trains the policy using the collected states, actions and rewards.
  #
  # Args:
  #   goal: goal coordinates, array with 2 entries.
  #   gamma (optional): discount factor.
  #   beta (optional): probability of slipping in the transition model.

  # Randomize starting position.
  cur_pos <- goal
  while(all(cur_pos == goal))
    cur_pos <- c(sample(1:H, size = 1), sample(1:W, size = 1))

  states <- NULL
  actions <- NULL

  steps <- 0 # To avoid getting stuck and/or training on unnecessarily long episodes.
  while(steps < 20){
    steps <- steps+1

    # Follow policy and execute action.
    action <- DeepPolicy(cur_pos[1], cur_pos[2], goal[1], goal[2])
    new_pos <- transition_model(cur_pos[1], cur_pos[2], action, beta)

    # Store states and actions.
    states <- c(states, cur_pos)
    actions <- c(actions, action)
    cur_pos <- new_pos

    if(all(new_pos == goal)){
      # Train network.
      DeepPolicy_train(states, actions, goal, gamma)
      break
    }
  }
}

```

Environment D:

2.6. Environment D. In this task, we will use eight goal positions for training and, then, validate the learned policy on the remaining eight possible goal positions. The training and validation goal positions are stored in the lists `train_goals` and `val_goals` in the code in the file `RL Lab2.R`. You are requested to run the code provided, which runs the REINFORCE algorithm for 5000 episodes with $\beta = 0$ and $\gamma = 0.95$.³ Each training episode uses a random goal position from `train_goals`. The initial position for the episode is also chosen at random. When training is completed, the code validates the learned policy for the goal positions in `val_goals`. This is done by with the help of the function `vis_prob`, which shows the grid, goal position and learned policy. Note that each non-terminal tile has four values. These represent the action probabilities associated to the tile (state). Note also that each non-terminal tile has an arrow. This indicates the action with the largest probability for the tile (ties are broken at random). Finally, answer the following questions:

- Has the agent learned a good policy? Why / Why not ?
- Could you have used the Q-learning algorithm to solve this task ?

Answer:

Given a training goals: `train_goals <- (4,1), (4,3), (3,1), (3,4), (2,1), (2,2), (1,2), (1,3)`
and validation goals: `val_goals <- (4,2), (4,4), (3,2), (3,3), (2,3), (2,4), (1,1), (1,4)`

The agent could find its way to the validation goals because it was trained on a training goals that are similar to the validation goals. The training goals gave the agent a good view of the environment from the first, second, third, and fourth rows, and that's why the agent was able to find its way to the validation goals with high probability. In other words, the agent was able to generalize the policy for similar states.

We can't use Q-Learning because it doesn't have the feature of generalization, for instance, if we set a different goal than the one Q-Learning was trained to find, then Q-learning won't be able to reach that goal. Reinforce has the generalization advantage which enables it to determine the action for states that it has not visited before given a similar state. So, if Reinforce visited a state that's similar to the state in interest, then it will give a similarly accurate result for that state.

```
# Environment D (training with random goal positions)

H <- 4
W <- 4

# Define the neural network (two hidden layers of 32 units each).
model <- keras_model_sequential()
model %>%
  layer_dense(units = 32, input_shape = c(4), activation = 'relu') %>%
  layer_dense(units = 32, activation = 'relu') %>%
  layer_dense(units = 4, activation = 'softmax')

compile(model, loss = "categorical_crossentropy", optimizer = optimizer_sgd(lr=0.001))

initial_weights <- get_weights(model)

train_goals <- list(c(4,1), c(4,3), c(3,1), c(3,4), c(2,1), c(2,2), c(1,2), c(1,3))
val_goals <- list(c(4,2), c(4,4), c(3,2), c(3,3), c(2,3), c(2,4), c(1,1), c(1,4))

show_validation <- function(epochs){

  for(goal in val_goals)
    vis_prob(goal, epochs)
```

```

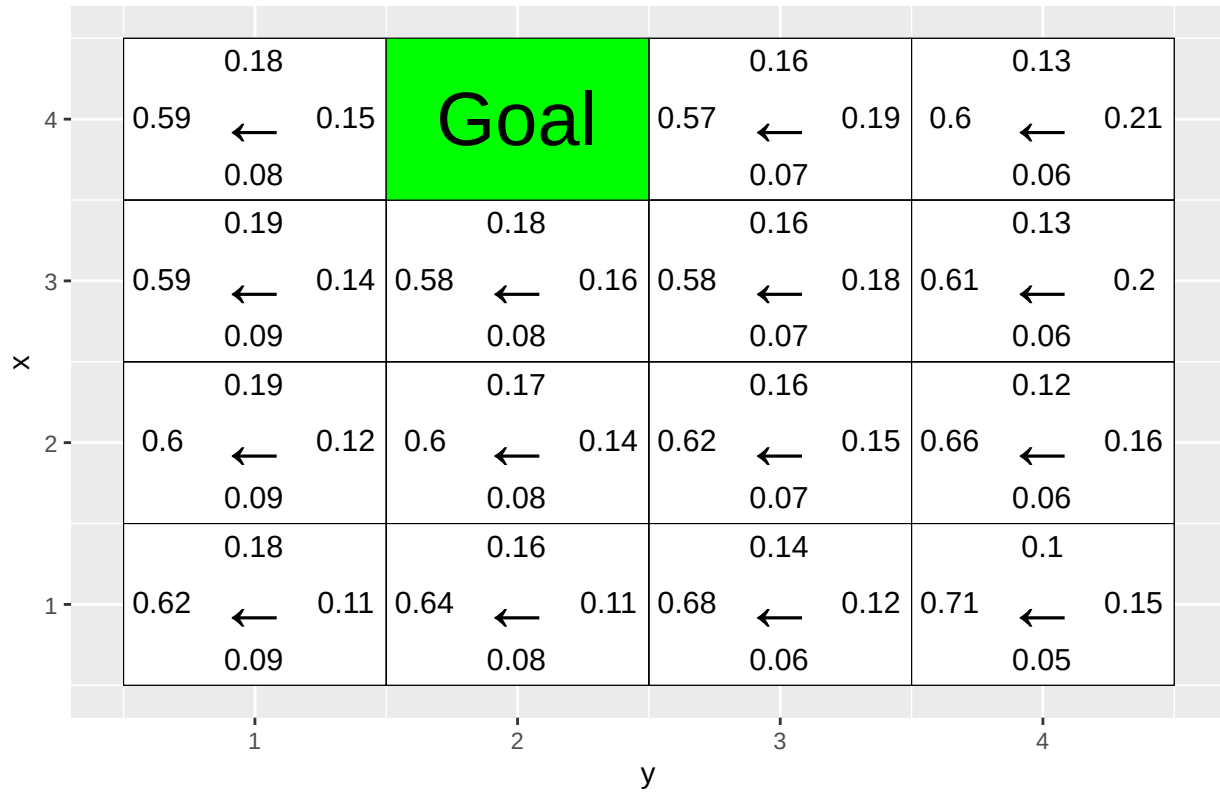
}

set_weights(model,initial_weights)

show_validation(0)

```


Action probabilities after 0 episodes




4	<div> <div>0.11</div> <div>0.73 ← 0.12</div> <div>0.04</div> </div>	<div> <div>0.1</div> <div>0.72 ← 0.14</div> <div>0.04</div> </div>	<div> <div>0.09</div> <div>0.71 ← 0.16</div> <div>0.03</div> </div>	Goal
3	<div> <div>0.11</div> <div>0.75 ← 0.1</div> <div>0.04</div> </div>	<div> <div>0.1</div> <div>0.74 ← 0.12</div> <div>0.04</div> </div>	<div> <div>0.09</div> <div>0.75 ← 0.13</div> <div>0.03</div> </div>	<div> <div>0.07</div> <div>0.77 ← 0.14</div> <div>0.03</div> </div>
2	<div> <div>0.1</div> <div>0.77 ← 0.08</div> <div>0.04</div> </div>	<div> <div>0.09</div> <div>0.78 ← 0.09</div> <div>0.04</div> </div>	<div> <div>0.07</div> <div>0.79 ← 0.1</div> <div>0.03</div> </div>	<div> <div>0.07</div> <div>0.8 ← 0.11</div> <div>0.02</div> </div>
1	<div> <div>0.09</div> <div>0.8 ← 0.07</div> <div>0.04</div> </div>	<div> <div>0.08</div> <div>0.82 ← 0.07</div> <div>0.03</div> </div>	<div> <div>0.07</div> <div>0.83 ← 0.08</div> <div>0.03</div> </div>	<div> <div>0.06</div> <div>0.83 ← 0.09</div> <div>0.02</div> </div>
	1	2	3	4

4	<div> <div>0.18</div> <div>0.54 ← 0.19</div> <div>0.09</div> </div>	<div> <div>0.17</div> <div>0.53 ← 0.21</div> <div>0.08</div> </div>	<div> <div>0.15</div> <div>0.53 ← 0.24</div> <div>0.08</div> </div>	<div> <div>0.12</div> <div>0.55 ← 0.27</div> <div>0.06</div> </div>
3	<div> <div>0.19</div> <div>0.53 ← 0.18</div> <div>0.1</div> </div>	Goal	<div> <div>0.15</div> <div>0.54 ← 0.23</div> <div>0.08</div> </div>	<div> <div>0.12</div> <div>0.56 ← 0.25</div> <div>0.06</div> </div>
2	<div> <div>0.19</div> <div>0.54 ← 0.16</div> <div>0.1</div> </div>	<div> <div>0.17</div> <div>0.55 ← 0.18</div> <div>0.09</div> </div>	<div> <div>0.15</div> <div>0.58 ← 0.2</div> <div>0.08</div> </div>	<div> <div>0.11</div> <div>0.61 ← 0.22</div> <div>0.06</div> </div>
1	<div> <div>0.19</div> <div>0.58 ← 0.14</div> <div>0.1</div> </div>	<div> <div>0.16</div> <div>0.6 ← 0.15</div> <div>0.09</div> </div>	<div> <div>0.14</div> <div>0.62 ← 0.17</div> <div>0.07</div> </div>	<div> <div>0.1</div> <div>0.65 ← 0.2</div> <div>0.05</div> </div>
	1	2	3	4


Action probabilities after 0 episodes




Action probabilities after 0 episodes



Action probabilities after 0 episodes



Action probabilities after 0 episodes



Action probabilities after 0 episodes

4		0.09 0.53 ← 0.33 0.04	0.08 0.49 ← 0.39 0.04	0.08 0.52 ← 0.37 0.03	0.08 0.56 ← 0.33 0.02				
3		0.09 0.54 ← 0.33 0.04	0.08 0.55 ← 0.34 0.04	0.08 0.58 ← 0.31 0.03	0.08 0.63 ← 0.26 0.02				
2		0.09 0.59 ← 0.28 0.04	0.08 0.61 ← 0.28 0.04	0.08 0.64 ← 0.25 0.03	0.08 0.68 ← 0.22 0.02				
1		0.08 0.64 ← 0.23 0.04	0.08 0.66 ← 0.22 0.04	0.08 0.69 ← 0.2 0.03	Goal				
		1	2	3	4				
		y							
x									

```
for(i in 1:5000){
  if(i%10==0) cat("episode",i,"\n")
  goal <- sample(train_goals, size = 1)
  reinforce_episode(unlist(goal))
}
```

```
## episode 10
## episode 20
## episode 30
## episode 40
## episode 50
## episode 60
## episode 70
## episode 80
## episode 90
## episode 100
## episode 110
## episode 120
## episode 130
## episode 140
## episode 150
## episode 160
## episode 170
## episode 180
## episode 190
## episode 200
## episode 210
```

episode 220
episode 230
episode 240
episode 250
episode 260
episode 270
episode 280
episode 290
episode 300
episode 310
episode 320
episode 330
episode 340
episode 350
episode 360
episode 370
episode 380
episode 390
episode 400
episode 410
episode 420
episode 430
episode 440
episode 450
episode 460
episode 470
episode 480
episode 490
episode 500
episode 510
episode 520
episode 530
episode 540
episode 550
episode 560
episode 570
episode 580
episode 590
episode 600
episode 610
episode 620
episode 630
episode 640
episode 650
episode 660
episode 670
episode 680
episode 690
episode 700
episode 710
episode 720
episode 730
episode 740
episode 750

episode 760
episode 770
episode 780
episode 790
episode 800
episode 810
episode 820
episode 830
episode 840
episode 850
episode 860
episode 870
episode 880
episode 890
episode 900
episode 910
episode 920
episode 930
episode 940
episode 950
episode 960
episode 970
episode 980
episode 990
episode 1000
episode 1010
episode 1020
episode 1030
episode 1040
episode 1050
episode 1060
episode 1070
episode 1080
episode 1090
episode 1100
episode 1110
episode 1120
episode 1130
episode 1140
episode 1150
episode 1160
episode 1170
episode 1180
episode 1190
episode 1200
episode 1210
episode 1220
episode 1230
episode 1240
episode 1250
episode 1260
episode 1270
episode 1280
episode 1290

episode 1300
episode 1310
episode 1320
episode 1330
episode 1340
episode 1350
episode 1360
episode 1370
episode 1380
episode 1390
episode 1400
episode 1410
episode 1420
episode 1430
episode 1440
episode 1450
episode 1460
episode 1470
episode 1480
episode 1490
episode 1500
episode 1510
episode 1520
episode 1530
episode 1540
episode 1550
episode 1560
episode 1570
episode 1580
episode 1590
episode 1600
episode 1610
episode 1620
episode 1630
episode 1640
episode 1650
episode 1660
episode 1670
episode 1680
episode 1690
episode 1700
episode 1710
episode 1720
episode 1730
episode 1740
episode 1750
episode 1760
episode 1770
episode 1780
episode 1790
episode 1800
episode 1810
episode 1820
episode 1830

episode 1840
episode 1850
episode 1860
episode 1870
episode 1880
episode 1890
episode 1900
episode 1910
episode 1920
episode 1930
episode 1940
episode 1950
episode 1960
episode 1970
episode 1980
episode 1990
episode 2000
episode 2010
episode 2020
episode 2030
episode 2040
episode 2050
episode 2060
episode 2070
episode 2080
episode 2090
episode 2100
episode 2110
episode 2120
episode 2130
episode 2140
episode 2150
episode 2160
episode 2170
episode 2180
episode 2190
episode 2200
episode 2210
episode 2220
episode 2230
episode 2240
episode 2250
episode 2260
episode 2270
episode 2280
episode 2290
episode 2300
episode 2310
episode 2320
episode 2330
episode 2340
episode 2350
episode 2360
episode 2370

episode 2380
episode 2390
episode 2400
episode 2410
episode 2420
episode 2430
episode 2440
episode 2450
episode 2460
episode 2470
episode 2480
episode 2490
episode 2500
episode 2510
episode 2520
episode 2530
episode 2540
episode 2550
episode 2560
episode 2570
episode 2580
episode 2590
episode 2600
episode 2610
episode 2620
episode 2630
episode 2640
episode 2650
episode 2660
episode 2670
episode 2680
episode 2690
episode 2700
episode 2710
episode 2720
episode 2730
episode 2740
episode 2750
episode 2760
episode 2770
episode 2780
episode 2790
episode 2800
episode 2810
episode 2820
episode 2830
episode 2840
episode 2850
episode 2860
episode 2870
episode 2880
episode 2890
episode 2900
episode 2910

episode 2920
episode 2930
episode 2940
episode 2950
episode 2960
episode 2970
episode 2980
episode 2990
episode 3000
episode 3010
episode 3020
episode 3030
episode 3040
episode 3050
episode 3060
episode 3070
episode 3080
episode 3090
episode 3100
episode 3110
episode 3120
episode 3130
episode 3140
episode 3150
episode 3160
episode 3170
episode 3180
episode 3190
episode 3200
episode 3210
episode 3220
episode 3230
episode 3240
episode 3250
episode 3260
episode 3270
episode 3280
episode 3290
episode 3300
episode 3310
episode 3320
episode 3330
episode 3340
episode 3350
episode 3360
episode 3370
episode 3380
episode 3390
episode 3400
episode 3410
episode 3420
episode 3430
episode 3440
episode 3450


episode 3460
episode 3470
episode 3480
episode 3490
episode 3500
episode 3510
episode 3520
episode 3530
episode 3540
episode 3550
episode 3560
episode 3570
episode 3580
episode 3590
episode 3600
episode 3610
episode 3620
episode 3630
episode 3640
episode 3650
episode 3660
episode 3670
episode 3680
episode 3690
episode 3700
episode 3710
episode 3720
episode 3730
episode 3740
episode 3750
episode 3760
episode 3770
episode 3780
episode 3790
episode 3800
episode 3810
episode 3820
episode 3830
episode 3840
episode 3850
episode 3860
episode 3870
episode 3880
episode 3890
episode 3900
episode 3910
episode 3920
episode 3930
episode 3940
episode 3950
episode 3960
episode 3970
episode 3980
episode 3990

episode 4000
episode 4010
episode 4020
episode 4030
episode 4040
episode 4050
episode 4060
episode 4070
episode 4080
episode 4090
episode 4100
episode 4110
episode 4120
episode 4130
episode 4140
episode 4150
episode 4160
episode 4170
episode 4180
episode 4190
episode 4200
episode 4210
episode 4220
episode 4230
episode 4240
episode 4250
episode 4260
episode 4270
episode 4280
episode 4290
episode 4300
episode 4310
episode 4320
episode 4330
episode 4340
episode 4350
episode 4360
episode 4370
episode 4380
episode 4390
episode 4400
episode 4410
episode 4420
episode 4430
episode 4440
episode 4450
episode 4460
episode 4470
episode 4480
episode 4490
episode 4500
episode 4510
episode 4520
episode 4530


```
## episode 4540
## episode 4550
## episode 4560
## episode 4570
## episode 4580
## episode 4590
## episode 4600
## episode 4610
## episode 4620
## episode 4630
## episode 4640
## episode 4650
## episode 4660
## episode 4670
## episode 4680
## episode 4690
## episode 4700
## episode 4710
## episode 4720
## episode 4730
## episode 4740
## episode 4750
## episode 4760
## episode 4770
## episode 4780
## episode 4790
## episode 4800
## episode 4810
## episode 4820
## episode 4830
## episode 4840
## episode 4850
## episode 4860
## episode 4870
## episode 4880
## episode 4890
## episode 4900
## episode 4910
## episode 4920
## episode 4930
## episode 4940
## episode 4950
## episode 4960
## episode 4970
## episode 4980
## episode 4990
## episode 5000
```

```
show_validation(5000)
```



Action probabilities after 5000 episodes




Action probabilities after 5000 episodes




Action probabilities after 5000 episodes




Action probabilities after 5000 episodes




Action probabilities after 5000 episodes




Action probabilities after 5000 episodes



Action probabilities after 5000 episodes



Action probabilities after 5000 episodes



Environment E:

2.7. Environment E. You are now asked to repeat the previous experiments but this time the goals for training are all from the top row of the grid. The validation goals are three positions from the rows below. To solve this task, simply run the code provided in the file RL Lab2.R and answer the following questions:

- Has the agent learned a good policy? Why / Why not ?
- If the results obtained for environments D and E differ, explain why.

Answer:

The agent couldn't find its way to the validation goals from most of the states because the training was done on the first row and the agent is unable to generalize over a very different goals from the ones it was trained on. The closer the goal to the first row, the more accurately the generalization works. Even though, the agent was able to reach the goal from some of the states, but in most cases (the majority of states), it was unable to find its way to the validation goal.

The result of environment E is different from that of D because the training goals are different and the validation goals are different. In D, the agent was trained on random goals from the environment with good exploration of different states, however, in E the agent was trained to reach goals on the first row only, which has made it difficult for it to reach validation goals that are far from the first row (unsimilar to the training goals).

```
# Environment E (training with top row goal positions)
```


```
train_goals <- list(c(4,1), c(4,2), c(4,3), c(4,4))
```

```
val_goals <- list(c(3,4), c(2,3), c(1,1))
```


```
set_weights(model,initial_weights)
```

```
show_validation(0)
```

Action probabilities after 0 episodes



Action probabilities after 0 episodes



4	<div> <div>0.2</div> <div>0.39 ← 0.25</div> <div>0.16</div> </div>	<div> <div>0.19</div> <div>0.39 ← 0.28</div> <div>0.14</div> </div>	<div> <div>0.18</div> <div>0.44 ← 0.27</div> <div>0.11</div> </div>	<div> <div>0.15</div> <div>0.51 ← 0.25</div> <div>0.08</div> </div>
3	<div> <div>0.21</div> <div>0.36 ← 0.27</div> <div>0.16</div> </div>	<div> <div>0.19</div> <div>0.39 ← 0.28</div> <div>0.14</div> </div>	<div> <div>0.18</div> <div>0.44 ← 0.28</div> <div>0.1</div> </div>	<div> <div>0.15</div> <div>0.53 ← 0.25</div> <div>0.07</div> </div>
2	<div> <div>0.21</div> <div>0.36 ← 0.27</div> <div>0.16</div> </div>	<div> <div>0.19</div> <div>0.39 ← 0.29</div> <div>0.13</div> </div>	<div> <div>0.17</div> <div>0.45 ← 0.29</div> <div>0.09</div> </div>	<div> <div>0.14</div> <div>0.51 ← 0.28</div> <div>0.06</div> </div>
1	<div> <div>Goal</div> </div>	<div> <div>0.19</div> <div>0.41 ← 0.28</div> <div>0.13</div> </div>	<div> <div>0.16</div> <div>0.46 ← 0.29</div> <div>0.09</div> </div>	<div> <div>0.14</div> <div>0.53 ← 0.28</div> <div>0.06</div> </div>
	1	2	3	4

```
## episode 10
## episode 20
## episode 30
## episode 40
## episode 50
## episode 60
## episode 70
## episode 80
## episode 90
## episode 100
## episode 110
## episode 120
## episode 130
## episode 140
## episode 150
## episode 160
## episode 170
## episode 180
## episode 190
## episode 200
## episode 210
```

episode 220
episode 230
episode 240
episode 250
episode 260
episode 270
episode 280
episode 290
episode 300
episode 310
episode 320
episode 330
episode 340
episode 350
episode 360
episode 370
episode 380
episode 390
episode 400
episode 410
episode 420
episode 430
episode 440
episode 450
episode 460
episode 470
episode 480
episode 490
episode 500
episode 510
episode 520
episode 530
episode 540
episode 550
episode 560
episode 570
episode 580
episode 590
episode 600
episode 610
episode 620
episode 630
episode 640
episode 650
episode 660
episode 670
episode 680
episode 690
episode 700
episode 710
episode 720
episode 730
episode 740
episode 750

episode 760
episode 770
episode 780
episode 790
episode 800
episode 810
episode 820
episode 830
episode 840
episode 850
episode 860
episode 870
episode 880
episode 890
episode 900
episode 910
episode 920
episode 930
episode 940
episode 950
episode 960
episode 970
episode 980
episode 990
episode 1000
episode 1010
episode 1020
episode 1030
episode 1040
episode 1050
episode 1060
episode 1070
episode 1080
episode 1090
episode 1100
episode 1110
episode 1120
episode 1130
episode 1140
episode 1150
episode 1160
episode 1170
episode 1180
episode 1190
episode 1200
episode 1210
episode 1220
episode 1230
episode 1240
episode 1250
episode 1260
episode 1270
episode 1280
episode 1290

episode 1300
episode 1310
episode 1320
episode 1330
episode 1340
episode 1350
episode 1360
episode 1370
episode 1380
episode 1390
episode 1400
episode 1410
episode 1420
episode 1430
episode 1440
episode 1450
episode 1460
episode 1470
episode 1480
episode 1490
episode 1500
episode 1510
episode 1520
episode 1530
episode 1540
episode 1550
episode 1560
episode 1570
episode 1580
episode 1590
episode 1600
episode 1610
episode 1620
episode 1630
episode 1640
episode 1650
episode 1660
episode 1670
episode 1680
episode 1690
episode 1700
episode 1710
episode 1720
episode 1730
episode 1740
episode 1750
episode 1760
episode 1770
episode 1780
episode 1790
episode 1800
episode 1810
episode 1820
episode 1830

episode 1840
episode 1850
episode 1860
episode 1870
episode 1880
episode 1890
episode 1900
episode 1910
episode 1920
episode 1930
episode 1940
episode 1950
episode 1960
episode 1970
episode 1980
episode 1990
episode 2000
episode 2010
episode 2020
episode 2030
episode 2040
episode 2050
episode 2060
episode 2070
episode 2080
episode 2090
episode 2100
episode 2110
episode 2120
episode 2130
episode 2140
episode 2150
episode 2160
episode 2170
episode 2180
episode 2190
episode 2200
episode 2210
episode 2220
episode 2230
episode 2240
episode 2250
episode 2260
episode 2270
episode 2280
episode 2290
episode 2300
episode 2310
episode 2320
episode 2330
episode 2340
episode 2350
episode 2360
episode 2370

episode 2380
episode 2390
episode 2400
episode 2410
episode 2420
episode 2430
episode 2440
episode 2450
episode 2460
episode 2470
episode 2480
episode 2490
episode 2500
episode 2510
episode 2520
episode 2530
episode 2540
episode 2550
episode 2560
episode 2570
episode 2580
episode 2590
episode 2600
episode 2610
episode 2620
episode 2630
episode 2640
episode 2650
episode 2660
episode 2670
episode 2680
episode 2690
episode 2700
episode 2710
episode 2720
episode 2730
episode 2740
episode 2750
episode 2760
episode 2770
episode 2780
episode 2790
episode 2800
episode 2810
episode 2820
episode 2830
episode 2840
episode 2850
episode 2860
episode 2870
episode 2880
episode 2890
episode 2900
episode 2910

episode 2920
episode 2930
episode 2940
episode 2950
episode 2960
episode 2970
episode 2980
episode 2990
episode 3000
episode 3010
episode 3020
episode 3030
episode 3040
episode 3050
episode 3060
episode 3070
episode 3080
episode 3090
episode 3100
episode 3110
episode 3120
episode 3130
episode 3140
episode 3150
episode 3160
episode 3170
episode 3180
episode 3190
episode 3200
episode 3210
episode 3220
episode 3230
episode 3240
episode 3250
episode 3260
episode 3270
episode 3280
episode 3290
episode 3300
episode 3310
episode 3320
episode 3330
episode 3340
episode 3350
episode 3360
episode 3370
episode 3380
episode 3390
episode 3400
episode 3410
episode 3420
episode 3430
episode 3440
episode 3450


episode 3460
episode 3470
episode 3480
episode 3490
episode 3500
episode 3510
episode 3520
episode 3530
episode 3540
episode 3550
episode 3560
episode 3570
episode 3580
episode 3590
episode 3600
episode 3610
episode 3620
episode 3630
episode 3640
episode 3650
episode 3660
episode 3670
episode 3680
episode 3690
episode 3700
episode 3710
episode 3720
episode 3730
episode 3740
episode 3750
episode 3760
episode 3770
episode 3780
episode 3790
episode 3800
episode 3810
episode 3820
episode 3830
episode 3840
episode 3850
episode 3860
episode 3870
episode 3880
episode 3890
episode 3900
episode 3910
episode 3920
episode 3930
episode 3940
episode 3950
episode 3960
episode 3970
episode 3980
episode 3990

episode 4000
episode 4010
episode 4020
episode 4030
episode 4040
episode 4050
episode 4060
episode 4070
episode 4080
episode 4090
episode 4100
episode 4110
episode 4120
episode 4130
episode 4140
episode 4150
episode 4160
episode 4170
episode 4180
episode 4190
episode 4200
episode 4210
episode 4220
episode 4230
episode 4240
episode 4250
episode 4260
episode 4270
episode 4280
episode 4290
episode 4300
episode 4310
episode 4320
episode 4330
episode 4340
episode 4350
episode 4360
episode 4370
episode 4380
episode 4390
episode 4400
episode 4410
episode 4420
episode 4430
episode 4440
episode 4450
episode 4460
episode 4470
episode 4480
episode 4490
episode 4500
episode 4510
episode 4520
episode 4530


```
## episode 4540
## episode 4550
## episode 4560
## episode 4570
## episode 4580
## episode 4590
## episode 4600
## episode 4610
## episode 4620
## episode 4630
## episode 4640
## episode 4650
## episode 4660
## episode 4670
## episode 4680
## episode 4690
## episode 4700
## episode 4710
## episode 4720
## episode 4730
## episode 4740
## episode 4750
## episode 4760
## episode 4770
## episode 4780
## episode 4790
## episode 4800
## episode 4810
## episode 4820
## episode 4830
## episode 4840
## episode 4850
## episode 4860
## episode 4870
## episode 4880
## episode 4890
## episode 4900
## episode 4910
## episode 4920
## episode 4930
## episode 4940
## episode 4950
## episode 4960
## episode 4970
## episode 4980
## episode 4990
## episode 5000
```

```
show_validation(5000)
```



Action probabilities after 5000 episodes



Action probabilities after 5000 episodes



Action probabilities after 5000 episodes



References

Course Documents
<https://stackoverflow.com/>

Appendix

```
RNGversion('3.5.1')
knitr::opts_chunk$set(echo = TRUE)
library(HMM)
library(entropy)
set.seed(12345)
library(ggplot2)

# If you do not see four arrows in line 16, then do the following:
# File/Reopen with Encoding/UTF-8

arrows <- c("↑", "→", "↓", "←")
action_deltas <- list(c(1,0), # up
                     c(0,1), # right
                     c(-1,0), # down
                     c(0,-1)) # left

vis_environment <- function(iterations=0, epsilon = 0.5, alpha = 0.1, gamma = 0.95, beta = 0){
```

```

# Visualize an environment with rewards.
# Q-values for all actions are displayed on the edges of each tile.
# The (greedy) policy for each state is also displayed.
#
# Args:
#   iterations, epsilon, alpha, gamma, beta (optional): for the figure title.
#   reward_map (global variable): a HxW array containing the reward given at each state.
#   q_table (global variable): a HxWx4 array containing Q-values for each state-action pair.
#   H, W (global variables): environment dimensions.

df <- expand.grid(x=1:H,y=1:W)
foo <- mapply(function(x,y) ifelse(reward_map[x,y] == 0,q_table[x,y,1],NA),df$x,df$y)
df$val1 <- as.vector(round(foo, 2))
foo <- mapply(function(x,y) ifelse(reward_map[x,y] == 0,q_table[x,y,2],NA),df$x,df$y)
df$val2 <- as.vector(round(foo, 2))
foo <- mapply(function(x,y) ifelse(reward_map[x,y] == 0,q_table[x,y,3],NA),df$x,df$y)
df$val3 <- as.vector(round(foo, 2))
foo <- mapply(function(x,y) ifelse(reward_map[x,y] == 0,q_table[x,y,4],NA),df$x,df$y)
df$val4 <- as.vector(round(foo, 2))
foo <- mapply(function(x,y)
  ifelse(reward_map[x,y] == 0,arrows[GreedyPolicy(x,y)],reward_map[x,y]),df$x,df$y)
df$val5 <- as.vector(foo)
foo <- mapply(function(x,y) ifelse(reward_map[x,y] == 0,max(q_table[x,y,]),
  ifelse(reward_map[x,y]<0,NA,reward_map[x,y])),df$x,df$y)
df$val6 <- as.vector(foo)

print(ggplot(df,aes(x = y,y = x)) +
  scale_fill_gradient(low = "white", high = "green", na.value = "red", name = "") +
  geom_tile(aes(fill=val6)) +
  geom_text(aes(label = val1),size = 4,nudge_y = .35,na.rm = TRUE) +
  geom_text(aes(label = val2),size = 4,nudge_x = .35,na.rm = TRUE) +
  geom_text(aes(label = val3),size = 4,nudge_y = -.35,na.rm = TRUE) +
  geom_text(aes(label = val4),size = 4,nudge_x = -.35,na.rm = TRUE) +
  geom_text(aes(label = val5),size = 10) +
  geom_tile(fill = 'transparent', colour = 'black') +
  ggtitle(paste("Q-table after ",iterations," iterations\n",
    "(epsilon = ",epsilon,", alpha = ",alpha,"gamma = ",gamma,", beta = ",beta,")")) +
  theme(plot.title = element_text(hjust = 0.5)) +
  scale_x_continuous(breaks = c(1:W),labels = c(1:W)) +
  scale_y_continuous(breaks = c(1:H),labels = c(1:H)))
}

GreedyPolicy <- function(x, y){

  # Get a greedy action for state (x,y) from q_table.
  #
  # Args:
  #   x, y: state coordinates.
  #   q_table (global variable): a HxWx4 array containing Q-values for each state-action pair.
  #
  # Returns:

```

```

# An action, i.e. integer in {1,2,3,4}.

# Your code here.
rewards = q_table[x,y,1:4]
action = 0
max_reward = which(rewards == max(rewards))

if (length(max_reward)>1){
  action = sample(max_reward,1)
}
else
{
  action = max_reward
}
return(action)
}

EpsilonGreedyPolicy <- function(x, y, epsilon){

  # Get an epsilon-greedy action for state (x,y) from q_table.
  #
  # Args:
  #   x, y: state coordinates.
  #   epsilon: probability of acting greedily.
  #
  # Returns:
  #   An action, i.e. integer in {1,2,3,4}.

  # Your code here.
  actions = c(1:4)
  action = 0
  rewards = q_table[x,y,1:4]

  max_reward = which(rewards == max(rewards))

  if (length(max_reward)>1){
    action = sample(max_reward,1)
  }
  else
  {
    action = max_reward
  }

  if (1-epsilon < runif(1)){
    return(sample(actions,1))
  }
  else{
    return(action)
  }
}

transition_model <- function(x, y, action, beta){

```

```

# Computes the new state after given action is taken. The agent will follow the action
# with probability (1-beta) and slip to the right or left with probability beta/2 each.
#
# Args:
#   x, y: state coordinates.
#   action: which action the agent takes (in {1,2,3,4}).
#   beta: probability of the agent slipping to the side when trying to move.
#   H, W (global variables): environment dimensions.
#
# Returns:
#   The new state after the action has been taken.

delta <- sample(-1:1, size = 1, prob = c(0.5*beta,1-beta,0.5*beta))
final_action <- ((action + delta + 3) %% 4) + 1
foo <- c(x,y) + unlist(action_deltas[final_action])
foo <- pmax(c(1,1),pmin(foo,c(H,W)))

return (foo)
}
q_learning <- function(start_state, epsilon = 0.5, alpha = 0.1, gamma = 0.95,
                        beta = 0){

# Perform one episode of Q-learning. The agent should move around in the
# environment using the given transition model and update the Q-table.
# The episode ends when the agent reaches a terminal state.
#
# Args:
#   start_state: array with two entries, describing the starting position of the agent.
#   epsilon (optional): probability of acting greedily.
#   alpha (optional): learning rate.
#   gamma (optional): discount factor.
#   beta (optional): slipping factor.
#   reward_map (global variable): a HxW array containing the reward given at each state.
#   q_table (global variable): a HxWx4 array containing Q-values for each state-action pair.
#
# Returns:
#   reward: reward received in the episode.
#   correction: sum of the temporal difference correction terms over the episode.
#   q_table (global variable): Recall that R passes arguments by value. So, q_table being
#   a global variable can be modified with the superassignment operator <<-.

# Your code here.
reward = 0
episode_correction = 0

repeat{
  # Follow policy, execute action, get reward.

  #Taking an action
  old_x = start_state[1]
  old_y = start_state[2]

  action= EpsilonGreedyPolicy(old_x, old_y, epsilon)

```

```

#Computing the new state after the action is taken
new_state = transition_model(old_x, old_y, action, beta)
new_x = new_state[1]
new_y = new_state[2]

#Computing the current reward
current_reward = reward_map[new_x, new_y]

# Q-table update.
#Getting a greedy action
maximum_action = GreedyPolicy(new_x,new_y)

#Calculating the final reward
reward = reward+current_reward

#Calculating the epsilon_correction
epsilon_correction = current_reward + gamma*q_table[new_x,new_y,maximum_action]-q_table[old_x,old_y,action]

q_table[old_x,old_y,action] <- q_table[old_x,old_y,action]+ alpha*epsilon_correction

start_state = new_state
if(reward!=0)
  # End episode.
  return (c(reward,epsilon_correction))
}

}

# Environment A (learning)

H <- 5
W <- 7

reward_map <- matrix(0, nrow = H, ncol = W)
reward_map[3,6] <- 10
reward_map[2:4,3] <- -1

q_table <- array(0,dim = c(H,W,4))

vis_environment()

for(i in 1:10000){
  foo <- q_learning(start_state = c(3,1))

  if(any(i==c(10,100,1000,10000)))
    vis_environment(i)
}

# Environment B (the effect of epsilon and gamma)

H <- 7

```

```

W <- 8

reward_map <- matrix(0, nrow = H, ncol = W)
reward_map[1,] <- -1
reward_map[7,] <- -1
reward_map[4,5] <- 5
reward_map[4,8] <- 10

q_table <- array(0,dim = c(H,W,4))

vis_environment()

MovingAverage <- function(x, n){

  cx <- c(0,cumsum(x))
  rsum <- (cx[(n+1):length(cx)] - cx[1:(length(cx) - n)]) / n

  return (rsum)
}

for(j in c(0.5,0.75,0.95)){
  q_table <- array(0,dim = c(H,W,4))
  reward <- NULL
  correction <- NULL

  for(i in 1:30000){
    foo <- q_learning(gamma = j, start_state = c(4,1))
    reward <- c(reward,foo[1])
    correction <- c(correction,foo[2])
  }

  vis_environment(i, gamma = j)
  plot(MovingAverage(reward,100),type = "l")
  plot(MovingAverage(correction,100),type = "l")
}

for(j in c(0.5,0.75,0.95)){
  q_table <- array(0,dim = c(H,W,4))
  reward <- NULL
  correction <- NULL

  for(i in 1:30000){
    foo <- q_learning(epsilon = 0.1, gamma = j, start_state = c(4,1))
    reward <- c(reward,foo[1])
    correction <- c(correction,foo[2])
  }

  vis_environment(i, epsilon = 0.1, gamma = j)
  plot(MovingAverage(reward,100),type = "l")
  plot(MovingAverage(correction,100),type = "l")
}

# Environment C (the effect of beta).

```

```

H <- 3
W <- 6

reward_map <- matrix(0, nrow = H, ncol = W)
reward_map[1,2:5] <- -1
reward_map[1,6] <- 10

q_table <- array(0,dim = c(H,W,4))

vis_environment()

for(j in c(0,0.2,0.4,0.66)){
  q_table <- array(0,dim = c(H,W,4))

  for(i in 1:10000)
    foo <- q_learning(gamma = 0.6, beta = j, start_state = c(1,1))

  vis_environment(i, gamma = 0.6, beta = j)
}
library(keras)

# install.packages("ggplot2")
# install.packages("vctrs")
library(ggplot2)

# If you do not see four arrows in line 19, then do the following:
# File/Reopen with Encoding/UTF-8

arrows <- c("↑", "→", "↓", "←")
action_deltas <- list(c(1,0), # up
                     c(0,1), # right
                     c(-1,0), # down
                     c(0,-1)) # left

vis_prob <- function(goal, episodes = 0){

  # Visualize an environment with rewards.
  # Probabilities for all actions are displayed on the edges of each tile.
  # The (greedy) policy for each state is also displayed.
  #
  # Args:
  #   goal: goal coordinates, array with 2 entries.
  #   episodes, epsilon, alpha, gamma, beta (optional): for the figure title.
  #   H, W (global variables): environment dimensions.

  df <- expand.grid(x=1:H,y=1:W)
  dist <- array(data = NA, dim = c(H,W,4))
  class <- array(data = NA, dim = c(H,W))
  for(i in 1:H)
    for(j in 1:W){
      dist[i,j,] <- DeepPolicy_dist(i,j,goal[1],goal[2])
      foo <- which(dist[i,j,]==max(dist[i,j,]))
      class[i,j] <- ifelse(length(foo)>1,sample(foo, size = 1),foo)
    }
}

```



```

}

foo <- mapapply(function(x,y) ifelse(all(c(x,y) == goal),NA,dist[x,y,1]),df$x,df$y)
df$val1 <- as.vector(round(foo, 2))
foo <- mapapply(function(x,y) ifelse(all(c(x,y) == goal),NA,dist[x,y,2]),df$x,df$y)
df$val2 <- as.vector(round(foo, 2))
foo <- mapapply(function(x,y) ifelse(all(c(x,y) == goal),NA,dist[x,y,3]),df$x,df$y)
df$val3 <- as.vector(round(foo, 2))
foo <- mapapply(function(x,y) ifelse(all(c(x,y) == goal),NA,dist[x,y,4]),df$x,df$y)
df$val4 <- as.vector(round(foo, 2))
foo <- mapapply(function(x,y) ifelse(all(c(x,y) == goal),NA,class[x,y]),df$x,df$y)
df$val5 <- as.vector(arrows[foo])
foo <- mapapply(function(x,y) ifelse(all(c(x,y) == goal),"Goal",NA),df$x,df$y)
df$val6 <- as.vector(foo)

print(ggplot(df,aes(x = y,y = x)) +
  geom_tile(fill = 'white', colour = 'black') +
  scale_fill_manual(values = c('green')) +
  geom_tile(aes(fill=val6), show.legend = FALSE, colour = 'black') +
  geom_text(aes(label = val1),size = 4,nudge_y = .35,na.rm = TRUE) +
  geom_text(aes(label = val2),size = 4,nudge_x = .35,na.rm = TRUE) +
  geom_text(aes(label = val3),size = 4,nudge_y = -.35,na.rm = TRUE) +
  geom_text(aes(label = val4),size = 4,nudge_x = -.35,na.rm = TRUE) +
  geom_text(aes(label = val5),size = 10,na.rm = TRUE) +
  geom_text(aes(label = val6),size = 10,na.rm = TRUE) +
  ggtitle(paste("Action probabilities after ",episodes," episodes")) +
  theme(plot.title = element_text(hjust = 0.5)) +
  scale_x_continuous(breaks = c(1:W),labels = c(1:W)) +
  scale_y_continuous(breaks = c(1:H),labels = c(1:H)))
}

transition_model <- function(x, y, action, beta){

  # Computes the new state after given action is taken. The agent will follow the action
  # with probability (1-beta) and slip to the right or left with probability beta/2 each.
  #
  # Args:
  #   x, y: state coordinates.
  #   action: which action the agent takes (in {1,2,3,4}).
  #   beta: probability of the agent slipping to the side when trying to move.
  #   H, W (global variables): environment dimensions.
  #
  # Returns:
  #   The new state after the action has been taken.

  delta <- sample(-1:1, size = 1, prob = c(0.5*beta,1-beta,0.5*beta))
  final_action <- ((action + delta + 3) %% 4) + 1
  foo <- c(x,y) + unlist(action_deltas[final_action])
  foo <- pmax(c(1,1),pmin(foo,c(H,W)))

  return (foo)
}

DeepPolicy_dist <- function(x, y, goal_x, goal_y){

```

```

# Get distribution over actions for state (x,y) and goal (goal_x,goal_y) from the deep policy.
#
# Args:
#   x, y: state coordinates.
#   goal_x, goal_y: goal coordinates.
#   model (global variable): NN encoding the policy.
#
# Returns:
#   A distribution over actions.

foo <- matrix(data = c(x,y,goal_x,goal_y), nrow = 1)

# return (predict_proba(model, x = foo))
return (predict_on_batch(model, x = foo)) # Faster.
}

DeepPolicy <- function(x, y, goal_x, goal_y){

  # Get an action for state (x,y) and goal (goal_x,goal_y) from the deep policy.
  #
  # Args:
  #   x, y: state coordinates.
  #   goal_x, goal_y: goal coordinates.
  #   model (global variable): NN encoding the policy.
  #
  # Returns:
  #   An action, i.e. integer in {1,2,3,4}.

  foo <- DeepPolicy_dist(x,y,goal_x,goal_y)

  return (sample(1:4, size = 1, prob = foo))
}

DeepPolicy_train <- function(states, actions, goal, gamma){

  # Train the policy network on a rolled out trajectory.
  #
  # Args:
  #   states: array of states visited throughout the trajectory.
  #   actions: array of actions taken throughout the trajectory.
  #   goal: goal coordinates, array with 2 entries.
  #   gamma: discount factor.

  # Construct batch for training.
  inputs <- matrix(data = states, ncol = 2, byrow = TRUE)
  inputs <- cbind(inputs,rep(goal[1],nrow(inputs)))
  inputs <- cbind(inputs,rep(goal[2],nrow(inputs)))

  targets <- array(data = actions, dim = nrow(inputs))
  targets <- to_categorical(targets-1, num_classes = 4)

  # Sample weights. Reward of 5 for reaching the goal.

```

```

weights <- array(data = 5*(gamma^(nrow(inputs)-1)), dim = nrow(inputs))

# Train on batch. Note that this runs a SINGLE gradient update.
train_on_batch(model, x = inputs, y = targets, sample_weight = weights)
}

reinforce_episode <- function(goal, gamma = 0.95, beta = 0){

  # Rolls out a trajectory in the environment until the goal is reached.
  # Then trains the policy using the collected states, actions and rewards.
  #
  # Args:
  #   goal: goal coordinates, array with 2 entries.
  #   gamma (optional): discount factor.
  #   beta (optional): probability of slipping in the transition model.

  # Randomize starting position.
  cur_pos <- goal
  while(all(cur_pos == goal))
    cur_pos <- c(sample(1:H, size = 1), sample(1:W, size = 1))

  states <- NULL
  actions <- NULL

  steps <- 0 # To avoid getting stuck and/or training on unnecessarily long episodes.
  while(steps < 20){
    steps <- steps+1

    # Follow policy and execute action.
    action <- DeepPolicy(cur_pos[1], cur_pos[2], goal[1], goal[2])
    new_pos <- transition_model(cur_pos[1], cur_pos[2], action, beta)

    # Store states and actions.
    states <- c(states, cur_pos)
    actions <- c(actions, action)
    cur_pos <- new_pos

    if(all(new_pos == goal)){
      # Train network.
      DeepPolicy_train(states, actions, goal, gamma)
      break
    }
  }
}

# Environment D (training with random goal positions)

H <- 4
W <- 4

# Define the neural network (two hidden layers of 32 units each).

```

```

model <- keras_model_sequential()
model %>%
  layer_dense(units = 32, input_shape = c(4), activation = 'relu') %>%
  layer_dense(units = 32, activation = 'relu') %>%
  layer_dense(units = 4, activation = 'softmax')

compile(model, loss = "categorical_crossentropy", optimizer = optimizer_sgd(lr=0.001))

initial_weights <- get_weights(model)

train_goals <- list(c(4,1), c(4,3), c(3,1), c(3,4), c(2,1), c(2,2), c(1,2), c(1,3))
val_goals <- list(c(4,2), c(4,4), c(3,2), c(3,3), c(2,3), c(2,4), c(1,1), c(1,4))

show_validation <- function(epochs){
  for(goal in val_goals)
    vis_prob(goal, epochs)
}

set_weights(model,initial_weights)

show_validation(0)

for(i in 1:5000){
  if(i%10==0) cat("episode",i,"\n")
  goal <- sample(train_goals, size = 1)
  reinforce_episode(unlist(goal))
}

show_validation(5000)

# Environment E (training with top row goal positions)

train_goals <- list(c(4,1), c(4,2), c(4,3), c(4,4))
val_goals <- list(c(3,4), c(2,3), c(1,1))

set_weights(model,initial_weights)

show_validation(0)

for(i in 1:5000){
  if(i%10==0) cat("episode", i,"\n")
  goal <- sample(train_goals, size = 1)
  reinforce_episode(unlist(goal))
}

show_validation(5000)

```